

Quasi-Monte Carlo vs Monte Carlo

for High-Dimensional Integration on $[0, 1]^d$

Tanner Wagner

MATH 505: Numerical Analysis

University of New Mexico

December 1, 2025

Outline

- 1 Motivation and Goal
- 2 Background
- 3 Test Integrands and Setup
- 4 Pseudocode (High-Level Overview)
- 5 Numerical Results
- 6 Discussion and Takeaways
- 7 Implementation Notes

High-dimensional integrals

- Many applications involve integrals over $[0, 1]^d$:
 - expectations in probability and statistics,
 - option pricing in mathematical finance,
 - posterior averages in Bayesian inference.
- Classical tensor-product quadrature (Gaussian, Newton–Cotes):
 - work grows exponentially with d ,
 - the *curse of dimensionality*.
- Need methods whose cost scales more gently with dimension.

Goal of the project

- Compare **plain Monte Carlo (MC)** with **Sobol quasi-Monte Carlo (QMC)** for integrals

$$I(f, d) = \int_{[0,1]^d} f(x) dx, \quad d \in \{5, 10, 15, 20\}.$$

- Test on several smooth, high-dimensional integrands with known exact values.
- Metrics:
 - absolute error,
 - empirical convergence rate,
 - **time-to-accuracy**.
- Question: *When does QMC actually beat MC in practice?*

Random variables in this talk

- Think of a **random variable** as:
 - a rule that assigns a number to the outcome of an experiment,
 - e.g. roll a die → number $1, \dots, 6$.
- In our setting:
 - $X = (X_1, \dots, X_d)$ is a random point in $[0, 1]^d$,
 - you can picture it as throwing a dart uniformly at the unit cube.
- A function value $f(X)$ is then also a random variable (it depends on where the dart lands).

i.i.d. samples and a coin-flip analogy

- **i.i.d.** = independent and identically distributed:
 - each sample is drawn with the same distribution,
 - and no sample influences any other.
- Coin-flip analogy:
 - flip a fair coin N times,
 - the fraction of heads is a sample average of i.i.d. Bernoulli r.v.'s.
- We do the same thing with integrals:
 - draw X_1, \dots, X_N i.i.d. uniform on $[0, 1]^d$,
 - look at the sample average of $f(X_i)$.

Integral as an expectation

- Let $X = (X_1, \dots, X_d)$ be uniform on $[0, 1]^d$.
- “Uniform” here means every region of $[0, 1]^d$ with the same volume is equally likely.
- Then

$$I(f, d) = \int_{[0,1]^d} f(x) dx = \mathbb{E}[f(X)].$$

- Approximate $I(f, d)$ by a sample mean:

$$\hat{I}_N = \frac{1}{N} \sum_{i=1}^N f(X_i),$$

with X_1, \dots, X_N i.i.d. $\text{Unif}([0, 1]^d)$.

- This is the basic idea of Monte Carlo integration: “*average many samples to approximate the true average.*”

Why Monte Carlo works (high-level picture)

- **Law of Large Numbers (LLN):**
 - sample average \hat{I}_N converges to the true mean $\mathbb{E}[f(X)]$ as $N \rightarrow \infty$,
 - just like the fraction of heads $\rightarrow \frac{1}{2}$ for a fair coin.
- Intuition:
 - each new sample adds a little information about the true average,
 - as N grows, the “random wiggles” cancel out.
- **Central Limit Theorem (CLT), informally:**
 - the error $\hat{I}_N - \mathbb{E}[f(X)]$ behaves like a Gaussian with width $\sim 1/\sqrt{N}$,
 - so typical error size is about a constant times $1/\sqrt{N}$.

Monte Carlo: error scaling

- Assume $\text{Var}(f(X)) = \sigma^2 < \infty$.

- Then

$$\text{Var}(\hat{I}_N) = \text{Var}\left(\frac{1}{N} \sum_{i=1}^N f(X_i)\right) = \frac{\sigma^2}{N}.$$

- Standard deviation (typical error size):

$$\sqrt{\text{Var}(\hat{I}_N)} = \frac{\sigma}{\sqrt{N}}.$$

- **Key point:**

$$\text{MC error} = O(N^{-1/2}),$$

with an exponent independent of d (though σ may depend on d).

Quasi-Monte Carlo (QMC)

- Replace random points X_i by a *low-discrepancy sequence* $(x_i)_{i \geq 1} \subset [0, 1]^d$:

$$\widehat{I}_N^{\text{QMC}}(f, d) = \frac{1}{N} \sum_{i=1}^N f(x_i).$$

- Sobol sequence: base-2 digital net designed to fill $[0, 1]^d$ very uniformly.
- For functions of bounded variation (Hardy–Krause), Koksma–Hlawka inequality gives

$$|\widehat{I}_N^{\text{QMC}} - I(f, d)| \leq V_{\text{HK}}(f) D^*(x_1, \dots, x_N),$$

where D^* is the star-discrepancy.

- For Sobol and related sequences,

$$D^*(x_1, \dots, x_N) = O\left(\frac{(\log N)^d}{N}\right).$$

What is a Sobol sequence?

- A **Sobol sequence** is a *low-discrepancy sequence* of points in $[0, 1]^d$:
 - deterministic (not i.i.d. random),
 - designed to be **very evenly spread out** over the cube.
- Built in base 2 as a *digital sequence*:
 - in 1D, think of a base-2 van der Corput sequence (points that successively fill in the gaps),
 - in higher d , Sobol chooses coordinates so that all low-dimensional projections are well distributed.
- Key property: for N points,

$$D^*(x_1, \dots, x_N) = O((\log N)^d / N),$$

so the points cover $[0, 1]^d$ much more uniformly than i.i.d. samples.

MC vs QMC: heuristic comparison

- Ignoring $(\log N)^d$:

$$\text{MC: } O(N^{-1/2}), \quad \text{QMC: } O(N^{-1}).$$

- In practice, QMC often gives

much smaller error than MC for the same N

when

- f is smooth,
- the *effective dimension* is moderate.
- But constants and effective dimension matter a lot in high d .
- This is exactly what we explore numerically.

Error metrics

For each method (MC or QMC) we consider:

- **Absolute error**

$$E_N(f, d) = |\widehat{I}_N(f, d) - I(f, d)|.$$

- **Error vs sample size:**

- log-log plots of E_N vs N ,
- estimate slopes (empirical convergence rates).

- **Error vs wall-clock time:**

- time-to-accuracy plots E_N vs T_N ,
- compare which method reaches a target tolerance faster.

- For MC, we also look at variability over repeated runs.

Dimensions and sample sizes

- Dimensions:

$$d \in \{5, 10, 15, 20\}.$$

- Sample sizes:

$$N \in \{2^7, 2^8, \dots, 2^{15}\} = \{128, 256, \dots, 32768\}.$$

- For each (f, d, N) :

- MC: $R = 30$ independent runs,
- QMC: single Sobol net of size N (via `random_base2` in SciPy).

Separable test integrands

- We choose $f : [0, 1]^d \rightarrow \mathbb{R}$ with closed-form integrals.
- First three integrands are *separable*:

$$f(x) = \prod_{j=1}^d g_j(x_j).$$

Then

$$I(f, d) = \prod_{j=1}^d \int_0^1 g_j(x_j) dx_j.$$

- This makes exact values cheap to compute while still giving nontrivial high-dimensional behavior.

Integrand A: separable exponential

- Define

$$f_A(x) = \exp\left(-\sum_{j=1}^d x_j\right) = \prod_{j=1}^d e^{-x_j}.$$

- Exact integral:

$$I(f_A, d) = \prod_{j=1}^d (1 - e^{-1}) = (1 - e^{-1})^d.$$

- Smooth, bounded, symmetric in all coordinates.

Integrand B: rational with arctan closed form

- Define

$$f_B(x) = \prod_{j=1}^d \frac{1}{1+jx_j^2}.$$

- 1D factor:

$$\int_0^1 \frac{1}{1+jx^2} dx = \frac{1}{\sqrt{j}} \arctan(\sqrt{j}).$$

- So

$$I(f_B, d) = \prod_{j=1}^d \frac{1}{\sqrt{j}} \arctan(\sqrt{j}).$$

- Smooth, but anisotropic: different behavior in different coordinate directions.

Integrand C: Gaussian-type

- Define

$$f_C(x) = \exp\left(-\sum_{j=1}^d x_j^2\right) = \prod_{j=1}^d e^{-x_j^2}.$$

- Use the error function:

$$\int_0^1 e^{-x^2} dx = \frac{\sqrt{\pi}}{2} \operatorname{erf}(1).$$

- Hence

$$I(f_C, d) = \left(\frac{\sqrt{\pi}}{2} \operatorname{erf}(1)\right)^d.$$

- Smooth, strongly related to Gaussian integrals.

Integrand D: low effective dimension

- Define, for $d \geq 5$,

$$f_D(x) = \left(\frac{1}{5} \sum_{j=1}^5 x_j \right)^2.$$

- Depends only on first five coordinates \Rightarrow *low effective dimension*.
- Let $X_1, \dots, X_5 \sim \text{Unif}(0, 1)$ independent, $S = X_1 + \dots + X_5$.
- Then

$$I(f_D, d) = \mathbb{E} \left[\left(\frac{S}{5} \right)^2 \right] = \frac{1}{25} (\text{Var}(S) + (\mathbb{E} S)^2) = \frac{4}{15},$$

independent of d .

- Designed to be especially favorable to QMC.

Plain Monte Carlo: high-level algorithm

Goal

Approximate

$$I(f, d) = \int_{[0,1]^d} f(x) dx$$

using random samples.

Pseudocode

- ① Input: integrand f , dimension d , sample size N .
- ② Set $\text{sum} \leftarrow 0$.
- ③ For $i = 1, \dots, N$:
 - ① draw $X_i \sim \text{Unif}([0, 1]^d)$
(e.g. each coordinate $X_i[j]$ is a uniform random number in $[0, 1]$),
 - ② $\text{sum} \leftarrow \text{sum} + f(X_i)$.
- ④ Output estimate

$$\hat{I}_N^{\text{MC}} = \frac{\text{sum}}{N}.$$

Sobol Quasi-Monte Carlo: high-level algorithm

Idea

Use a deterministic low-discrepancy sequence instead of random points.

Pseudocode


- ① Input: integrand f , dimension d , sample size N .
- ② Construct a Sobol sequence generator in dimension d .
- ③ Set $\text{sum} \leftarrow 0$.
- ④ For $i = 1, \dots, N$:
 - ① take the next point x_i from the Sobol sequence in $[0, 1]^d$,
 - ② $\text{sum} \leftarrow \text{sum} + f(x_i)$.
- ⑤ Output estimate

$$\hat{I}_N^{\text{QMC}} = \frac{\text{sum}}{N}.$$

Here there is no randomness: once you fix “Sobol” and the dimension d , the sequence (x_i) is deterministic. The hope is that these points cover $[0, 1]^d$ so evenly that the error behaves like const/N .


Error vs N : moderate dimension ($d = 5$)

- Example: f_A in $d = 5$.




- QMC curve lies below MC mean error for moderate and large N .
- On log–log scale:

Error vs N : high dimension ($d = 20$)




f_A in $d = 20$



f_B in $d = 20$



- QMC still shows $\approx N^{-1}$ -like decay but with a *large constant*.
- MC errors start much smaller and stay below QMC for all N tested.

Error vs N : Gaussian-type f_C in $d = 20$



- For small N : MC error < QMC error (pre-asymptotic regime).
- As N grows: QMC decays faster, eventually catching up and slightly overtaking MC.

Error vs N : low effective dimension f_D




f_D in $d = 5$

- QMC clearly dominates MC in both $d = 5$ and $d = 20$.
- QMC error curves in $d = 5$ and $d = 20$ are very similar: effective dimension ≈ 5 .
- Textbook example of QMC's $O(N^{-1})$ -type behavior.

Error vs time: representative examples

- For each configuration, we measure wall-clock time T_N .
- Plots of E_N vs T_N mirror the E_N vs N behavior:
 - When QMC has smaller error at fixed N (e.g. f_A in $d = 5$, f_D), it also reaches a target tolerance faster.
 - When MC dominates in error (e.g. f_A and f_B in $d = 20$), MC is also more efficient in time.
- Cost per sample is similar for MC and QMC, so accuracy differences translate directly into time-to-accuracy.


Error vs time: illustrative examples



f_A in $d = 5$

QMC reaches target accuracy faster

Same cost per sample for MC and QMC \Rightarrow the lower curve at a given time is the more efficient method.



f_A in $d = 20$

MC more efficient over this range

Summary of findings

- **Moderate dimension ($d = 5$):**
 - QMC outperforms MC on all three separable integrands f_A, f_B, f_C .
 - Empirical slopes: QMC ≈ -1 , MC $\approx -1/2$.
- **High dimension ($d = 20$):**
 - For f_A and f_B (very small exact integrals), MC has smaller error than QMC for all N .
 - For Gaussian-type f_C , MC is better for small N , QMC catches up at larger N .
- **Low effective dimension (f_D):**
 - QMC clearly wins in both $d = 5$ and $d = 20$.
 - Error curves essentially independent of the nominal dimension.

Rates, constants, and effective dimension

- Asymptotic statements

$$\text{MC } O(N^{-1/2}), \quad \text{QMC } O(N^{-1})$$

hide important constants.

- For f_A and f_B in $d = 20$:
 - exact integrals are extremely small,
 - MC errors inherit this small scale via the variance,
 - QMC behaves more like C/N with $C \approx O(1)$.
- Effective dimension matters:
 - when most variation is in a few coordinates (like f_D), QMC works extremely well,
 - when variation is spread across many coordinates, the benefit can be muted or delayed.

Practical guidelines

- Use Sobol QMC when:
 - the integrand is smooth,
 - effective dimension is moderate or low,
 - you can afford moderately large N .
- Plain MC can be competitive or better when:
 - dimension is very high,
 - exact integrals / variances are extremely small,
 - you are restricted to relatively small N .
- Overall message:

“QMC is great, but problem structure really matters.”

Python implementation (Code/)

- `integrands.py`
 - definitions of $f_A, f_B, f_C, f_D,$
 - closed-form formulas for $I(f, d).$
- `mc_qmc.py`
 - plain MC estimator using NumPy random sampling,
 - Sobol QMC estimator using `scipy.stats.qmc.Sobol`.
- `run_experiments.py`
 - loops over $f, d, N,$
 - performs $R = 30$ MC runs and 1 QMC run,
 - writes results to `results_mc_qmc.csv`.
- `plot_results.py`
 - reads `results_mc_qmc.csv`,
 - generates the error-vs- N and error-vs-time plots.

References

- H. Niederreiter, *Random Number Generation and Quasi–Monte Carlo Methods*, SIAM, 1992.
- C. Lemieux, *Monte Carlo and Quasi–Monte Carlo Sampling*, Springer, 2009.
- A. Quarteroni, R. Sacco, and F. Saleri, *Numerical Mathematics*, 2nd ed., Springer, 2007.
- A. B. Owen, *Monte Carlo Theory, Methods and Examples*, 2013. Available online at <https://artowen.su.domains/mc/>.

Questions?